

# TANISHA PRITHA

Nagpur, India | [E-mail](#) | [LinkedIn](#) | [GitHub](#) | [LeetCode](#) | 7368825023

## PROFESSIONAL SUMMARY

**Fullstack Engineer** with hands-on experience across **React / Next.js frontends, FastAPI and Node-style backends**, and **LLM-powered systems**. Comfortable owning features end-to-end — UI, REST APIs, PostgreSQL, and Redis-backed job queues — with additional depth in agentic AI workflows, RAG pipelines, and multi-agent orchestration.

## TECHNICAL SKILLS

**Frontend:** React.js, Next.js, TypeScript

**Backend:** Python, FastAPI, Django, REST APIs, Async I/O, JWT, RBAC, Redis, Supabase, Background Tasks, ETL Pipelines

**Data & Infra:** PostgreSQL, Redis, Docker, AWS EC2, Vercel, Streamlit, Scikit-learn, XGBoost, PaddleOCR

**AI / GenAI (Bonus):** LLMs, Production RAG, Retrieval Engineering, Agentic & Multi-Agent Workflows (LangGraph), Model Context Protocol (MCP), Structured Output (Pydantic), Prompt Engineering, AI Evaluation (RAGAS), Observability & Tracing (LangSmith)

**AI Infra:** LiteLLM, Cohere Rerank, Pinecone, pgvector, FAISS, ChromaDB, BAAI/bge, all-MiniLM, OpenRouter, Deepgram STT

**Strengths:** Fullstack Delivery, Retrieval Quality, Cost Optimization, AI System Observability, Security-First Engineering

## EXPERIENCE

**AI Engineer Intern | AptlyHired (Startup)** | Nov 2025 – Dec 2025 | Remote, New Delhi

- **Built an end-to-end RAG pipeline** with multi-format ingestion (PDF, DOCX, TXT, PNG, JPG), PaddleOCR, semantic chunking, pgvector retrieval, and an Upload → Store → Chat workflow in Django.
- Improved retrieval accuracy through optimized chunk sizing, noise filtering, and prompt design; integrated into a web UI used by internal testers.

## PROJECTS

**Threadbase | Content Scheduling Platform** | Next.js, React, FastAPI, PostgreSQL, Redis, Supabase, Clerk

[GitHub](#) | [Live App](#)

- **Built an LLM-powered content engine** that generates and schedules social posts from user ideas and media; exposed a REST API (draft, schedule, publish) with Clerk auth and JWT-protected endpoints.
- **Designed a Redis-backed job queue** with FastAPI BackgroundTasks and Redis pub/sub for async scheduling; idempotency keys prevent duplicate jobs and TTL-based cache invalidation reduces redundant LLM calls.
- Tracked job lifecycle (enqueued, executing, delivered) in PostgreSQL for a full audit trail; Supabase handles media storage with signed URL access.

**Multi-Agent Real Estate Sales Automation Platform** | LangGraph, MCP, Redis Streams, Celery, LiteLLM, LangSmith

[GitHub](#) | [Live App](#)

- **Built a stateful multi-agent AI workflow platform** using LangGraph to automate lead qualification, property recommendations, follow-ups, and CRM updates with durable execution.
- Standardized AI tool calling via **Model Context Protocol (MCP)** servers for CRM, inventory, messaging, calendar, and knowledge integrations; routed calls through **LiteLLM** for multi-model fallback.
- Improved workflow reliability with event-driven orchestration (**Redis Streams, Celery**, checkpointing, retries, human-in-the-loop approval); instrumented with **LangSmith** tracing and a Streamlit ops dashboard for real-time monitoring.

**Multi-Agent Company Compliance Engine** | FastAPI, pgvector, LangSmith, Cohere Rerank, RAGAS

[GitHub](#) | [Live App](#)

- **Architected a Planner-Reasoner-Verifier multi-agent system** for forensic auditing of compliance documents; agent consensus across all three agents required before decisions are finalized, cutting hallucination risk in high-stakes outputs.
- **Built a production RAG pipeline** with semantic chunking, BAAI/bge embeddings, pgvector retrieval, and **Cohere Rerank** to cut top-K from 20 to 5; enforced **Pydantic structured output** (risk level, violations, confidence) on all LLM responses.

## EDUCATION

**B.Tech in Mechanical Engineering** | Government College of Engineering, Nagpur | Graduating: June, 2027

## CERTIFICATIONS

Introduction to Machine Learning: Art of the Possible (AWS) | HackerRank SQL (Intermediate) | Data Modeling for Amazon DocumentDB (AWS)